

## CONTRIBUTIONS

We introduce Deep Quality-Value (DQV) Learning, a novel Deep Reinforcement Learning (DRL) algorithm which learns significantly faster and better than Deep Q-Learning and Double Deep Q-Learning. DQV uses temporal-difference learning to train a Value neural network and uses this network for training a second Quality-value network that learns to estimate state-action values.

## PRELIMINARIES AND MOTIVATION

We consider a set of possible states,  $\mathcal{S}$ , and a set of possible actions  $\mathcal{A}$  that allow an agent to perform a state transition from state  $s_t$  at time  $t$  to  $s_{t+1}$  defined by a transition probability distribution  $p(s_{t+1}|s_t, a_t)$ . Associated to it there is an immediate reward function,  $\mathcal{R}(s_t, a_t, s_{t+1})$ . Actions are taken based on a policy  $\pi : s \rightarrow a$ . For each state we can calculate its Value  $V$  with respect to a discount factor  $\gamma$ :

$$V^\pi(s) = E \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k} \middle| s_t = s \right],$$

whereas values of state-action pairs,  $Q^\pi(s, a)$  can be maximized with a policy that satisfies

$$Q^\pi(s_t, a_t) = \sum_{s_{t+1} \in \mathcal{S}} p(s_{t+1}|s_t, a_t) \left( \mathcal{R}(s_t, a_t, s_{t+1}) + \gamma \max_{a_{t+1} \in \mathcal{A}} Q^\pi(s_{t+1}, a_{t+1}) \right).$$

**Quality-Value  $QV(\lambda)$  Learning** is an online tabular RL algorithm proposed in [1] which keeps track of both the  $Q$  function and the  $V$  function. First  $V$  is learned through temporal difference TD( $\lambda$ ) learning and the following update rule:

$$V(s) := V(s) + \alpha [r_t + \gamma V(s_{t+1}) - V(s_t)] e_t(s).$$

Then these estimates are used to learn the  $Q$  function with an update rule which is similar to one step Q-Learning.

$$Q(s_t, a_t) := Q(s_t, a_t) + \alpha [r_t + \gamma V(s_{t+1}) - Q(s_t, a_t)].$$

**Goal:** We aim to formulate these update rules as objective functions which can be minimized via gradient descent.

## DQV-LEARNING

We use two neural networks parametrized as  $\Phi$  and  $\theta$  respectively, and formulate  $QV(\lambda)$ 's update rules in Mean Squared Error terms. We can now minimize the  $V$  function with:

$$L_\Phi = E \left[ \underbrace{(r_t + \gamma V(s_{t+1}, \Phi) - V(s_t, \Phi))^2}_{y_t} \right]$$

and the  $Q$  function with:

$$L_\theta = E \left[ \underbrace{(r_t + \gamma V(s_{t+1}, \Phi) - Q(s_t, a_t, \theta))^2}_{y_t} \right].$$

Note that both neural networks use the same target  $y_t$  when learning. We additionally improve the stability of DQV with:

- **Experience Replay:** a memory buffer,  $D$ , of size  $N$ , which stores experiences as  $\langle s_t, a_t, r_t, s_{t+1} \rangle$  that we can uniformly sample from  $\langle s_t, a_t, r_t, s_{t+1} \rangle \sim U(D)$  for optimizing the neural networks.
- **Target Neural Network:** a separate neural network, parametrized as  $\Phi^-$  that is specifically designed for estimating the targets,  $y_t$ , that are necessary for computing the TD errors. This slightly modifies the original loss functions to:

$$L_\Phi = E \left[ (r_t + \gamma V(s_{t+1}, \Phi^-) - V(s_t, \Phi))^2 \right]$$

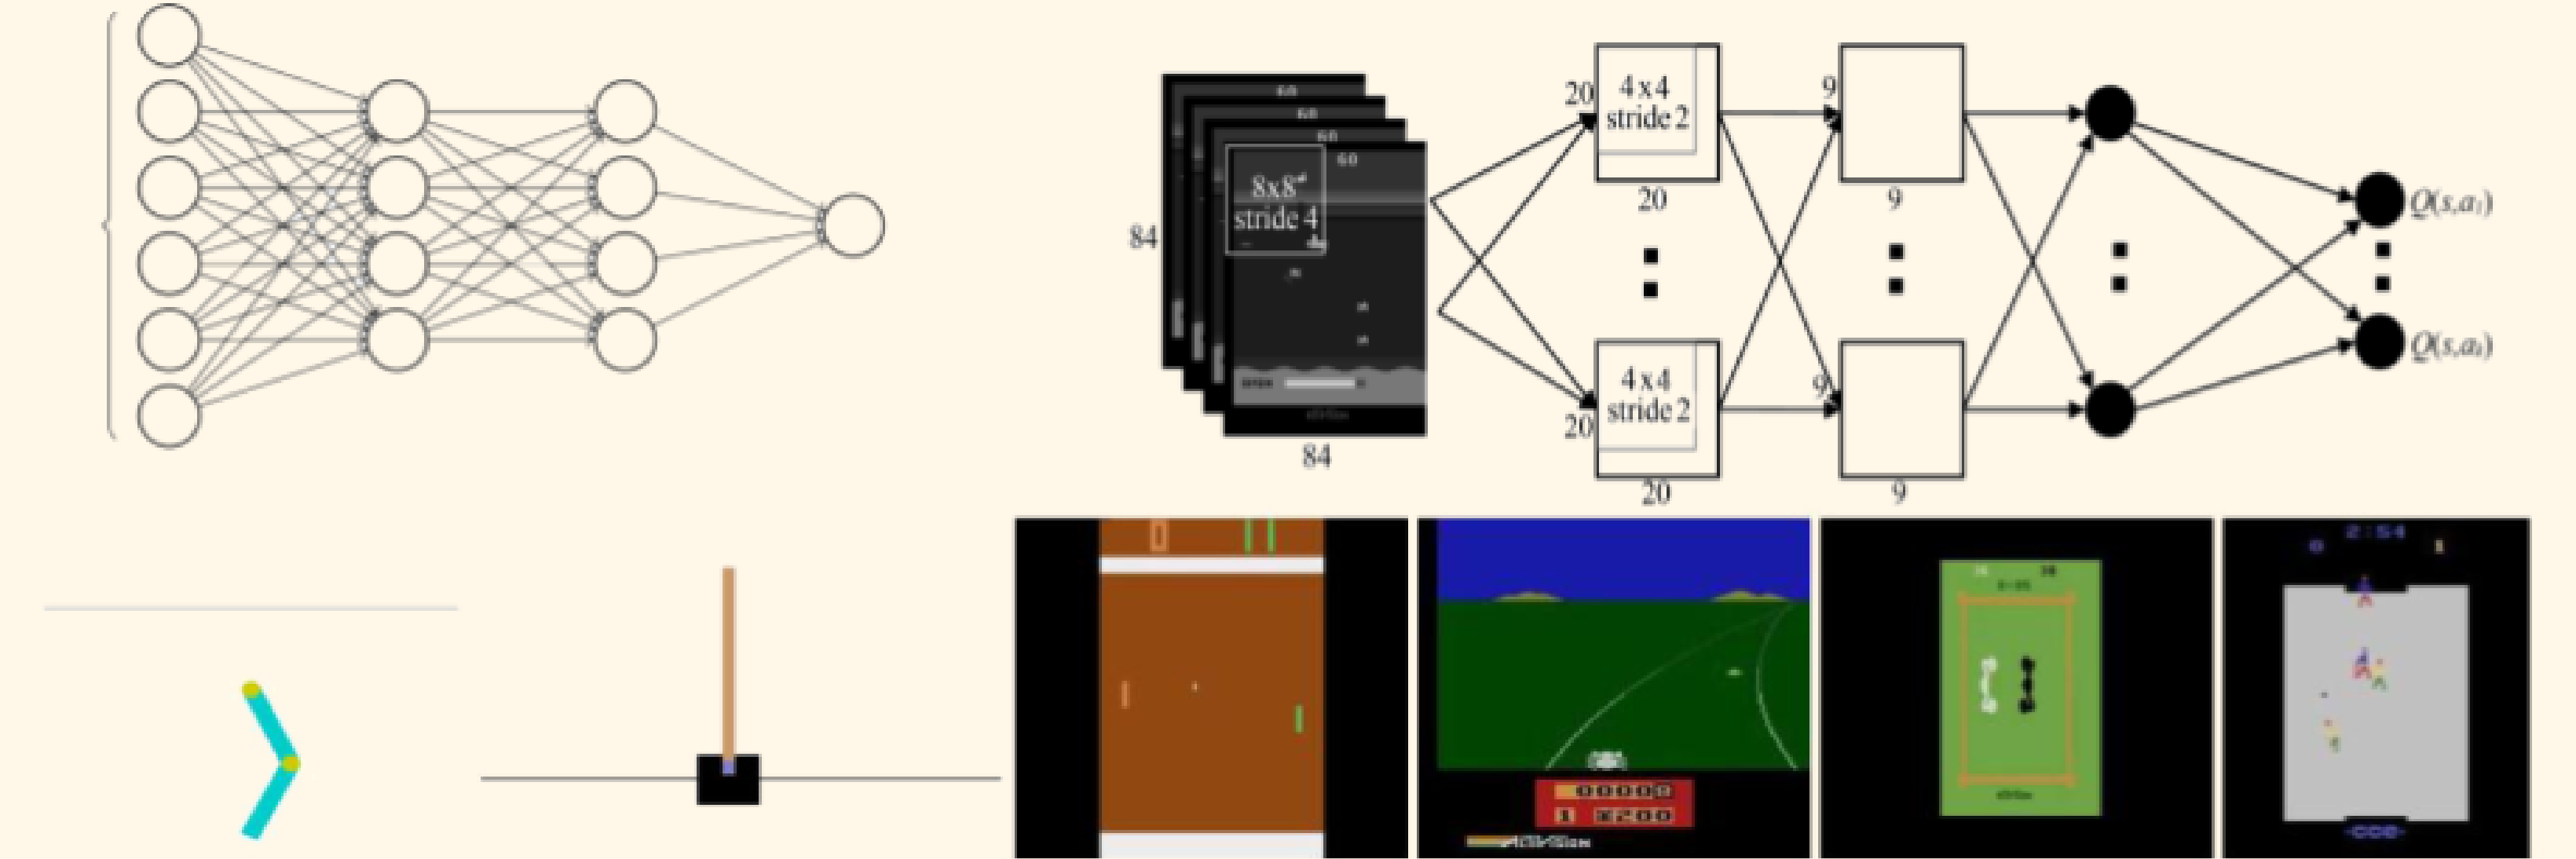
and

$$L_\theta = E \left[ (r_t + \gamma V(s_{t+1}, \Phi^-) - Q(s_t, a_t, \theta))^2 \right].$$

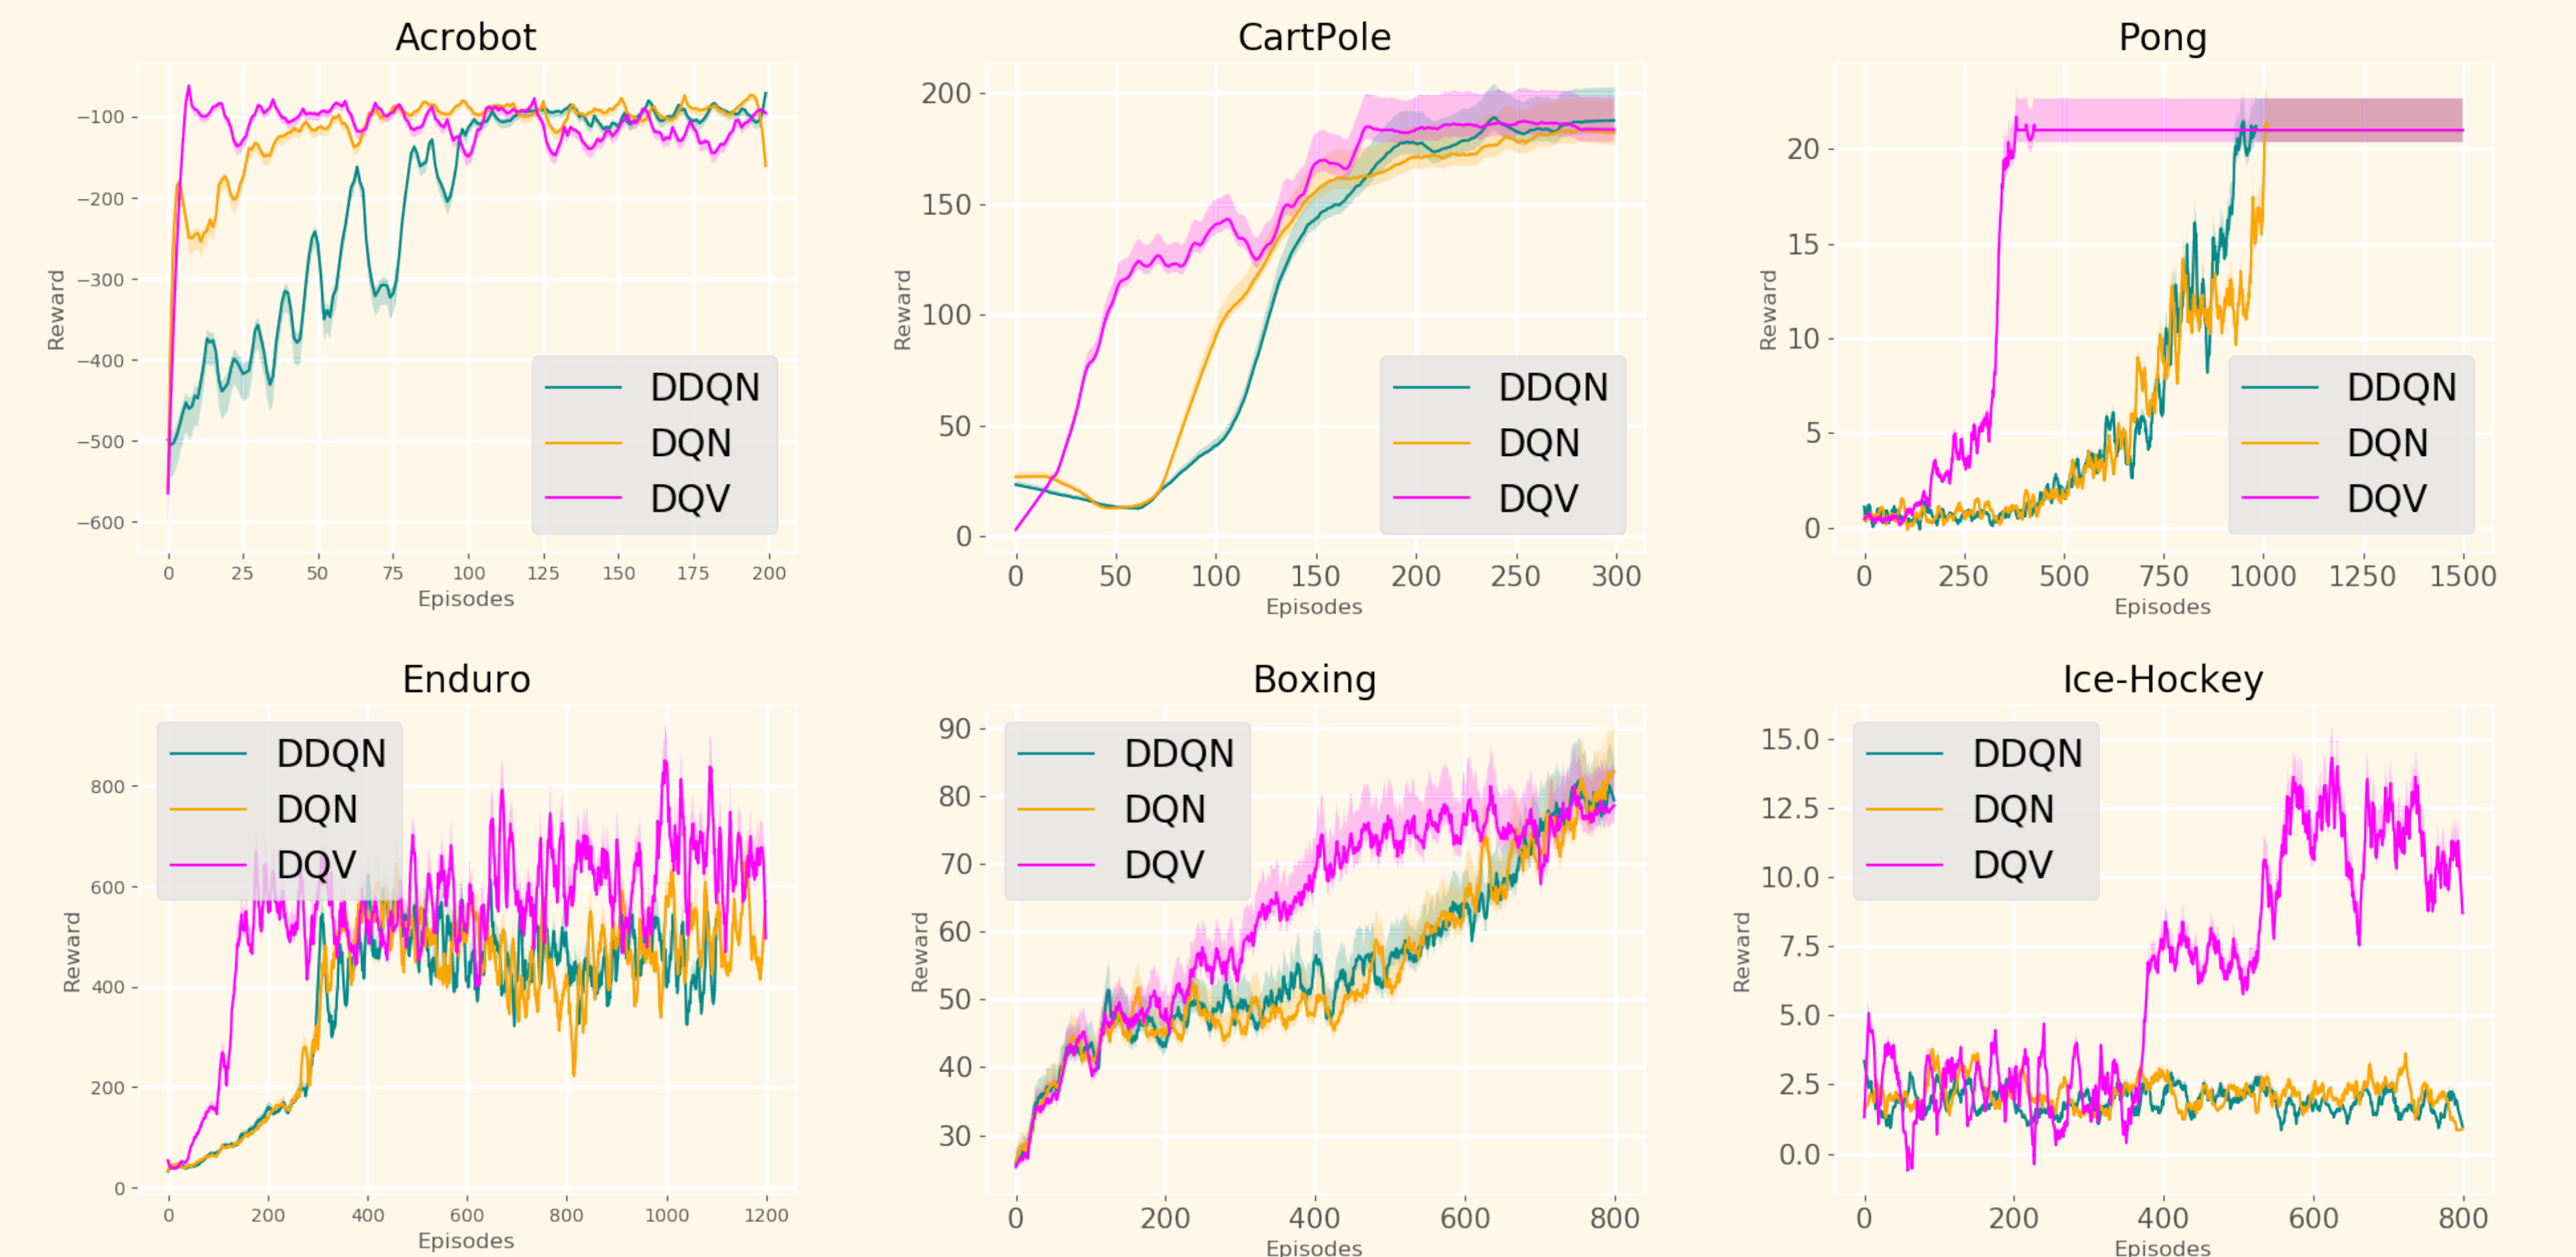
We update the Value-Target-Network  $\Phi^-$  with the weights of our original Value Network  $\Phi$  every 10,000 actions as defined by the hyperparameter  $c$ .

## EXPERIMENTAL SETUP

We use Multilayer Perceptrons for approximating the  $V$  and the  $Q$  function on two classic RL problems, and Deep Convolutional Neural Networks on four Atari games. The environments that have been used are from left to right: *Acrobot*, *Cartpole*, *Pong*, *Enduro*, *Boxing* and *Ice-Hockey*.



## RESULTS AND DISCUSSION



When used in combination with a Multilayer Perceptron on two classic RL problems (Acrobot and Cartpole) DQV learns significantly faster when compared with DQN [2] and DDQN [3]. Similarly, when used in combination with Deep Convolutional Neural Networks, Experience Replay and Target Neural Networks on four Atari games, DQV learns faster on the games *Pong* and *Boxing* while it also yields better results on the games *Enduro* and *Ice-Hockey*.

## ACKNOWLEDGEMENTS

Matthia Sabatelli acknowledges the financial support of the Federal Public Planning Service Science Policy, Belgium, within the BRAIN-be project.

## REFERENCES

- [1] Marco A. Wiering QV( $\lambda$ )-learning: A new on-policy reinforcement learning algorithm.
- [2] Mnih et al. Human-level control through deep reinforcement learning.
- [3] Van Hasselt et al. Deep Reinforcement Learning with Double Q-Learning.